

main

September 26, 2024

##

DA5401 Data Analytics Labarotary

###

Assignment 6 - Submitted by: DA24M011 - Nandhakishore C S

Task 1 [40 Points] Given the dataset (Amazon's MASSIVE dataset) - has data from 51 Languages in the latest version. For the given question, we are downloading the 27 files corresponding to languages which use the roman / latin script.

Importing Libraries

```
[1]: import os
import numpy as np
from datasets import load_dataset
from unicode import unicode # type: ignore

import unicodedata
from sklearn.naive_bayes import MultinomialNB
from sklearn.preprocessing import LabelEncoder
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay, accuracy_score
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
import pandas as pd
```

0.0.1 TASK 1

Defining the languages which use Latin Script and downloading the corresponding files from the main dataset.

```
[2]: # Define the locales we are interested in (i.e.) the languages which use Latin
    Script
locales = [
    'af-ZA', 'da-DK', 'de-DE', 'en-US', 'es-ES', 'fr-FR', 'fi-FI', 'hu-HU', \
    'is-IS', 'it-IT',
    'jv-ID', 'lv-LV', 'ms-MY', 'nb-NO', 'nl-NL', 'pl-PL', 'pt-PT', 'ro-RO', \
    'ru-RU', 'sl-SL',
```

```

    'sv-SE', 'sq-AL', 'sw-KE', 'tl-PH', 'tr-TR', 'vi-VN', 'cy-GB'
]

```

```

[3]: # from datasets import load_dataset - using Huggingface load dataset module /_
      ↪function
      df_raw = load_dataset("AmazonScience/massive", "all")

```

```

[4]: df_raw

```

```

[4]: DatasetDict({
    train: Dataset({
        features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
        num_rows: 587214
    })
    validation: Dataset({
        features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
        num_rows: 103683
    })
    test: Dataset({
        features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
        num_rows: 151674
    })
})

```

```

[5]: # Downloading and storing datafiles for 27 languages

output_dir = 'language_files'
os.makedirs(output_dir, exist_ok=True)

# Extract utterances for each locale
for locale in locales:
    file_path = os.path.join(output_dir, f'{locale}.txt')
    with open(file_path, 'w', encoding='utf-8') as f:
        for partition in ['train', 'validation', 'test']:
            for example in df_raw[partition]:
                if example['locale'] == locale:
                    utt = example['utt']
                    f.write(utt + '\n')

```

0.0.2 TASK 2

The dataset is split into three parttions:

1. Train
2. Validation

3. Test

From the training partition, we get the ‘utt’ column as the input features (the sentences with / without) accents and the corresponding country name in the ‘locale’ column as the label. Instead of using the tokens from the dataset, Tokens are generated using CountVectoriser from sklearn’s feature extraction module.

Using ‘filter’ function to get the data corresponding to the country names

```
[6]: df_filtered = df_raw.filter(lambda x: x['locale'] in locales)
df_filtered

[6]: DatasetDict({
  train: Dataset({
    features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
    num_rows: 310878
  })
  validation: Dataset({
    features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
    num_rows: 54891
  })
  test: Dataset({
    features: ['id', 'locale', 'partition', 'scenario', 'intent', 'utt',
'annot_utt', 'worker_id', 'slot_method', 'judgments'],
    num_rows: 80298
  })
})
```

Building Multinomial Naive Bayes Model - without removing accents

```
[7]: # No accent removal
df1_train = df_filtered['train']
df1_val = df_filtered['validation']
df1_test = df_filtered['test']

[8]: pipeline = Pipeline([
  ('vectoriser', CountVectorizer()),
  ('classifier', MultinomialNB())
])

pipeline.fit(df1_train['utt'], df1_train['locale'])
```

```
[8]: Pipeline(steps=[('vectoriser', CountVectorizer()),
  ('classifier', MultinomialNB())])
```

Model’s performance metrics for Train Partition

```
[9]: train_predictions = pipeline.predict(df1_train['utt'])
print("Training Data Performance:")
print(classification_report(df1_train['locale'], train_predictions))
print("Training Accuracy:", accuracy_score(df1_train['locale'],
↪train_predictions))
```

Training Data Performance:

	precision	recall	f1-score	support
af-ZA	0.98	0.98	0.98	11514
cy-GB	1.00	1.00	1.00	11514
da-DK	0.97	0.97	0.97	11514
de-DE	1.00	0.99	0.99	11514
en-US	0.96	0.99	0.98	11514
es-ES	0.99	0.99	0.99	11514
fi-FI	1.00	0.99	0.99	11514
fr-FR	0.99	0.99	0.99	11514
hu-HU	1.00	0.99	1.00	11514
is-IS	1.00	1.00	1.00	11514
it-IT	0.99	0.99	0.99	11514
jv-ID	0.99	0.99	0.99	11514
lv-LV	1.00	1.00	1.00	11514
ms-MY	0.99	0.99	0.99	11514
nb-NO	0.97	0.96	0.97	11514
nl-NL	0.99	0.98	0.98	11514
pl-PL	0.99	0.99	0.99	11514
pt-PT	0.99	0.99	0.99	11514
ro-RO	1.00	0.99	0.99	11514
ru-RU	1.00	1.00	1.00	11514
sl-SL	1.00	0.99	1.00	11514
sq-AL	1.00	0.99	1.00	11514
sv-SE	0.99	0.99	0.99	11514
sw-KE	1.00	1.00	1.00	11514
tl-PH	1.00	1.00	1.00	11514
tr-TR	1.00	0.99	1.00	11514
vi-VN	1.00	1.00	1.00	11514
accuracy			0.99	310878
macro avg	0.99	0.99	0.99	310878
weighted avg	0.99	0.99	0.99	310878

Training Accuracy: 0.9909160506693945

Model's performance metrics for Validation Partition

```
[10]: val_predictions = pipeline.predict(df1_val['utt'])
print("Validation Data Performance:")
print(classification_report(df1_val['locale'], val_predictions))
```

```
print("Validation Accuracy:", accuracy_score(df1_val['locale'],
↪val_predictions))
```

Validation Data Performance:

	precision	recall	f1-score	support
af-ZA	0.91	0.98	0.94	2033
cy-GB	1.00	0.99	0.99	2033
da-DK	0.94	0.96	0.95	2033
de-DE	1.00	0.98	0.99	2033
en-US	0.96	0.99	0.98	2033
es-ES	0.98	0.98	0.98	2033
fi-FI	1.00	0.98	0.99	2033
fr-FR	0.99	0.99	0.99	2033
hu-HU	1.00	0.98	0.99	2033
is-IS	1.00	0.99	0.99	2033
it-IT	0.98	0.99	0.99	2033
jv-ID	0.99	0.98	0.99	2033
lv-LV	1.00	0.99	0.99	2033
ms-MY	0.99	0.99	0.99	2033
nb-NO	0.96	0.94	0.95	2033
nl-NL	0.98	0.97	0.98	2033
pl-PL	0.99	0.98	0.98	2033
pt-PT	0.98	0.98	0.98	2033
ro-RO	1.00	0.99	0.99	2033
ru-RU	1.00	0.99	1.00	2033
sl-SL	1.00	0.99	0.99	2033
sq-AL	1.00	0.99	0.99	2033
sv-SE	0.97	0.98	0.97	2033
sw-KE	1.00	0.99	1.00	2033
tl-PH	0.99	0.99	0.99	2033
tr-TR	1.00	0.99	0.99	2033
vi-VN	1.00	1.00	1.00	2033
accuracy			0.98	54891
macro avg	0.98	0.98	0.98	54891
weighted avg	0.98	0.98	0.98	54891

Validation Accuracy: 0.9838589204058953

Model's performance metrics for Test Partition

```
[11]: test_predictions = pipeline.predict(df1_test['utt'])
print("Testing Data Performance:")
print(classification_report(df1_test['locale'], test_predictions))
print("Testing Accuracy:", accuracy_score(df1_test['locale'], test_predictions))
```

Testing Data Performance:

	precision	recall	f1-score	support
af-ZA	0.89	0.98	0.94	2974
cy-GB	1.00	0.99	1.00	2974
da-DK	0.94	0.95	0.95	2974
de-DE	0.99	0.99	0.99	2974
en-US	0.94	0.99	0.97	2974
es-ES	0.98	0.98	0.98	2974
fi-FI	0.99	0.98	0.99	2974
fr-FR	0.99	0.99	0.99	2974
hu-HU	1.00	0.98	0.99	2974
is-IS	1.00	0.99	0.99	2974
it-IT	0.98	0.99	0.99	2974
jv-ID	0.99	0.98	0.99	2974
lv-LV	0.99	0.99	0.99	2974
ms-MY	0.99	0.99	0.99	2974
nb-NO	0.96	0.93	0.95	2974
nl-NL	0.99	0.97	0.98	2974
pl-PL	1.00	0.98	0.99	2974
pt-PT	0.98	0.98	0.98	2974
ro-RO	1.00	0.99	0.99	2974
ru-RU	1.00	0.99	1.00	2974
sl-SL	1.00	0.99	0.99	2974
sq-AL	1.00	0.99	0.99	2974
sv-SE	0.99	0.97	0.98	2974
sw-KE	1.00	0.99	1.00	2974
tl-PH	0.99	0.99	0.99	2974
tr-TR	0.99	0.99	0.99	2974
vi-VN	1.00	1.00	1.00	2974
accuracy			0.98	80298
macro avg	0.98	0.98	0.98	80298
weighted avg	0.98	0.98	0.98	80298

Testing Accuracy: 0.9833868838576303

Now, building a MultinomialNB model with accent removal from the data. For the “unicodedata” library is used.

```
[28]: # ACCENT REMOVAL

df2 = df_filtered
# Function to deaccent characters
def deaccent(text):
    return ''.join(c for c in unicodedata.normalize('NFD', text) if unicodedata.
        category(c) != 'Mn')

# Create a directory to store the files
```

```
# Iterate through the dataset and save sentences to respective files
for locale in locales:
    for partition in ['train', 'validation', 'test']:
        for item in df2[partition]:
            if item['locale'] == locale:
                sentence = item['utt']
                sentence = deaccent(sentence)
```

```
[29]: df2_train = df2['train']
      df2_val = df2['validation']
      df2_test = df2['test']
```

Building Multinomial Naive Bayes Model - with accents

```
[30]: pipeline = Pipeline([
      ('vectoriser', CountVectorizer()),
      ('classifier', MultinomialNB())
    ])

    pipeline.fit(df2_train['utt'], df2_train['locale'])
```

```
[30]: Pipeline(steps=[('vectoriser', CountVectorizer()),
                      ('classifier', MultinomialNB())])
```

Model's performance metrics for Train Partition

```
[31]: train_predictions = pipeline.predict(df2_train['utt'])
      print("Training Data Performance:")
      print(classification_report(df2_train['locale'], train_predictions))
      print("Training Accuracy:", accuracy_score(df2_train['locale'],
      ↪train_predictions))
```

Training Data Performance:

	precision	recall	f1-score	support
af-ZA	0.98	0.98	0.98	11514
cy-GB	1.00	1.00	1.00	11514
da-DK	0.97	0.97	0.97	11514
de-DE	1.00	0.99	0.99	11514
en-US	0.96	0.99	0.98	11514
es-ES	0.99	0.99	0.99	11514
fi-FI	1.00	0.99	0.99	11514
fr-FR	0.99	0.99	0.99	11514
hu-HU	1.00	0.99	1.00	11514
is-IS	1.00	1.00	1.00	11514
it-IT	0.99	0.99	0.99	11514
jv-ID	0.99	0.99	0.99	11514

lv-LV	1.00	1.00	1.00	11514
ms-MY	0.99	0.99	0.99	11514
nb-NO	0.97	0.96	0.97	11514
nl-NL	0.99	0.98	0.98	11514
pl-PL	0.99	0.99	0.99	11514
pt-PT	0.99	0.99	0.99	11514
ro-RO	1.00	0.99	0.99	11514
ru-RU	1.00	1.00	1.00	11514
sl-SL	1.00	0.99	1.00	11514
sq-AL	1.00	0.99	1.00	11514
sv-SE	0.99	0.99	0.99	11514
sw-KE	1.00	1.00	1.00	11514
tl-PH	1.00	1.00	1.00	11514
tr-TR	1.00	0.99	1.00	11514
vi-VN	1.00	1.00	1.00	11514
accuracy			0.99	310878
macro avg	0.99	0.99	0.99	310878
weighted avg	0.99	0.99	0.99	310878

Training Accuracy: 0.9909160506693945

Model's performance metrics for Validation Partition

```
[32]: val_predictions = pipeline.predict(df2_val['utt'])
print("Validation Data Performance:")
print(classification_report(df2_val['locale'], val_predictions))
print("Validation Accuracy:", accuracy_score(df2_val['locale'],
↪val_predictions))
```

Validation Data Performance:

	precision	recall	f1-score	support
af-ZA	0.91	0.98	0.94	2033
cy-GB	1.00	0.99	0.99	2033
da-DK	0.94	0.96	0.95	2033
de-DE	1.00	0.98	0.99	2033
en-US	0.96	0.99	0.98	2033
es-ES	0.98	0.98	0.98	2033
fi-FI	1.00	0.98	0.99	2033
fr-FR	0.99	0.99	0.99	2033
hu-HU	1.00	0.98	0.99	2033
is-IS	1.00	0.99	0.99	2033
it-IT	0.98	0.99	0.99	2033
jv-ID	0.99	0.98	0.99	2033
lv-LV	1.00	0.99	0.99	2033
ms-MY	0.99	0.99	0.99	2033
nb-NO	0.96	0.94	0.95	2033

nl-NL	0.98	0.97	0.98	2033
pl-PL	0.99	0.98	0.98	2033
pt-PT	0.98	0.98	0.98	2033
ro-RO	1.00	0.99	0.99	2033
ru-RU	1.00	0.99	1.00	2033
sl-SL	1.00	0.99	0.99	2033
sq-AL	1.00	0.99	0.99	2033
sv-SE	0.97	0.98	0.97	2033
sw-KE	1.00	0.99	1.00	2033
tl-PH	0.99	0.99	0.99	2033
tr-TR	1.00	0.99	0.99	2033
vi-VN	1.00	1.00	1.00	2033
accuracy			0.98	54891
macro avg	0.98	0.98	0.98	54891
weighted avg	0.98	0.98	0.98	54891

Validation Accuracy: 0.9838589204058953

Model's performance metrics for Test Partition

```
[33]: test_predictions = pipeline.predict(df2_test['utt'])
print("Testing Data Performance:")
print(classification_report(df2_test['locale'], test_predictions))
print("Testing Accuracy:", accuracy_score(df2_test['locale'], test_predictions))
```

Testing Data Performance:

	precision	recall	f1-score	support
af-ZA	0.89	0.98	0.94	2974
cy-GB	1.00	0.99	1.00	2974
da-DK	0.94	0.95	0.95	2974
de-DE	0.99	0.99	0.99	2974
en-US	0.94	0.99	0.97	2974
es-ES	0.98	0.98	0.98	2974
fi-FI	0.99	0.98	0.99	2974
fr-FR	0.99	0.99	0.99	2974
hu-HU	1.00	0.98	0.99	2974
is-IS	1.00	0.99	0.99	2974
it-IT	0.98	0.99	0.99	2974
jv-ID	0.99	0.98	0.99	2974
lv-LV	0.99	0.99	0.99	2974
ms-MY	0.99	0.99	0.99	2974
nb-NO	0.96	0.93	0.95	2974
nl-NL	0.99	0.97	0.98	2974
pl-PL	1.00	0.98	0.99	2974
pt-PT	0.98	0.98	0.98	2974
ro-RO	1.00	0.99	0.99	2974

ru-RU	1.00	0.99	1.00	2974
sl-SL	1.00	0.99	0.99	2974
sq-AL	1.00	0.99	0.99	2974
sv-SE	0.99	0.97	0.98	2974
sw-KE	1.00	0.99	1.00	2974
tl-PH	0.99	0.99	0.99	2974
tr-TR	0.99	0.99	0.99	2974
vi-VN	1.00	1.00	1.00	2974
accuracy			0.98	80298
macro avg	0.98	0.98	0.98	80298
weighted avg	0.98	0.98	0.98	80298

Testing Accuracy: 0.9833868838576303

From the above results, We can see that keeping accents in the data is useful for language classification. Accuracy for model without removing accents is 99% and accuracy for model without accents is 98%. Note the fact that, the Precision for model with accents is higher than the precision for the model without accents.

TASK 3 Collapsing the chosen 27 languages into a 4 label dataset, where the labels are the continent they are spoken in. Building a Regularised Discriminant Analyser (RDA) over the above modified data with Linear Discriminant Analyser (LDA) and Quadratic Discriminant Analyser (QDA) with a parameter which balances the predictions of both.

```
[18]: continent_groups = {
    'af-ZA': 'Africa', 'sw-KE': 'Africa',
    'da-DK': 'Europe', 'de-DE': 'Europe', 'es-ES': 'Europe', 'fr-FR': 'Europe',
    ↪ 'fi-FI': 'Europe',
    'hu-HU': 'Europe', 'is-IS': 'Europe', 'it-IT': 'Europe', 'lv-LV': 'Europe',
    ↪ 'nb-NO': 'Europe',
    'nl-NL': 'Europe', 'pl-PL': 'Europe', 'pt-PT': 'Europe', 'ro-RO': 'Europe',
    ↪ 'ru-RU': 'Europe',
    'sl-SL': 'Europe', 'sv-SE': 'Europe', 'sq-AL': 'Europe', 'cy-GB': 'Europe',
    'jv-ID': 'Asia', 'ms-MY': 'Asia', 'tl-PH': 'Asia', 'tr-TR': 'Asia', 'vi-VN':
    ↪ 'Asia',
    'en-US': 'North America'
}
```

Creating a extra column in the dataset and adding the continent value to the corresponding language and then passing the new column as the label for the RDA Model

```
[19]: def assign_continent(row):
    locale = row['locale']
    return continent_groups.get(locale, 'NA')
```

Using 'lambda' functions in python to apply the above function to add the continent names.

```
[20]: train_data = df_filtered['train'].map(lambda row: {'continent':  

↳ assign_continent(row)})  

val_data = df_filtered['validation'].map(lambda row: {'continent':  

↳ assign_continent(row)})  

test_data = df_filtered['test'].map(lambda row: {'continent':  

↳ assign_continent(row)})
```

Class for Regularised Discriminant Analysis

```
[34]: from sklearn.base import BaseEstimator, ClassifierMixin # to get two  

↳ classifiers inside a single class  

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis,  

↳ QuadraticDiscriminantAnalysis  

from sklearn.metrics import accuracy_score  
  

class RegularizedDiscriminantAnalysis(BaseEstimator, ClassifierMixin):  

    __slots__ = '_lambda'  

    def __init__(self, _lambda=0.5):  

        self._lambda = _lambda  

        self.lda = LinearDiscriminantAnalysis()  

        self.qda = QuadraticDiscriminantAnalysis()  
  

    def fit(self, X, y):  

        self.lda.fit(X, y)  

        self.qda.fit(X, y)  

        return self  
  

    def predict(self, X):  

        lda_pred = self.lda.predict_proba(X)  

        qda_pred = self.qda.predict_proba(X)  

        combined_pred = ((1 - self._lambda) * lda_pred) + (self._lambda *  

↳ qda_pred)  

        return np.argmax(combined_pred, axis=1)
```

Using min frequency pruning (min_df = 10) and limiting the features to 250, the model is build and tuned for hyper parameter lambda

```
[55]: X_train_vec = CountVectorizer(max_features=400, min_df=10).  

↳ fit_transform(train_data['utt'])  

y_train = LabelEncoder().fit_transform(train_data['continent'])  
  

X_test_vec = CountVectorizer(max_features=400, min_df=10).  

↳ fit_transform(test_data['utt'])  

y_test = LabelEncoder().fit_transform(test_data['continent'])  
  

X_val_vec = CountVectorizer(max_features=400, min_df=10).  

↳ fit_transform(val_data['utt'])
```

```
y_val = LabelEncoder().fit_transform(val_data['continent'])
```

```
[56]: # Initialize and train the RDA model
rda = RegularizedDiscriminantAnalysis(_lambda=0)
rda.fit(X_train_vec.toarray(), y_train)
```

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are
collinear
  warnings.warn("Variables are collinear")
```

```
[56]: RegularizedDiscriminantAnalysis(_lambda=0)
```

```
[57]: y_test = LabelEncoder().fit_transform(y_test)
# Predict and evaluate the model
y_pred = rda.predict(X_test_vec.toarray())
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
```

```
Accuracy: 0.6566913248150639
```

```
[40]: # Define the parameter grid
parameter_grid = [0, 0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1]
best_accuracy = float('-inf')
best_parameter = 0

for i in range(0, len(parameter_grid)):
    model = RegularizedDiscriminantAnalysis(_lambda = parameter_grid[i]).
    ↪fit(X_train_vec.toarray(), y_train)
    y_pred = model.predict(X_val_vec.toarray())
    accuracy = accuracy_score(y_val, y_pred)
    if (accuracy > best_accuracy):
        best_parameter = parameter_grid[i]
        best_accuracy = accuracy
    print(f'Iteration{i} Done')
```

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are
collinear
  warnings.warn("Variables are collinear")
```

```
Iteration0 Done
```

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are
collinear
  warnings.warn("Variables are collinear")
```

```
Iteration1 Done
```

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration2 Done

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration3 Done

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration4 Done

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration5 Done

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration6 Done

```
/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-  
packages/sklearn/discriminant_analysis.py:926: UserWarning: Variables are  
collinear
```

```
warnings.warn("Variables are collinear")
```

Iteration7 Done

```
[41]: best_accuracy
```

```
[41]: 0.733599315006103
```

```
[42]: best_parameter
```

```
[42]: 0
```

After hyperparameter tuning, the best parameter is $\lambda = 0$. With the best parameter, the accuracy is 73.5% with $\text{max_df} = 10$ and $\text{max_features} = 400$ in counter vectoriser.