

DA24M011

September 15, 2024

#####

DA5401 Data Analytics Labarotary

#####

Assignment 5 - Submitted by: DA24M011 - Nandhakishore C S

Task 1 [40 Points] (From Question)

Let's consider the classification problem in <https://archive.ics.uci.edu/dataset/76/nursery> which is a 8-features, 3-classes dataset. It is mentioned in the link that the expected performance of over 90% accuracy (See Baseline Model Performance). Let's add the following model performance outcomes to the baselines, shall we?

1. Decision Tree (categorical features)
2. Decision Tree (categorical features in one-hot encoded form)
3. Logistic Regression with L1 regularization
4. K-Nearest Neighbors

You are expected to split the data into train, val & test. Use the val partition to tune the hyper-parameters such as (but not limited to) k of kNN, height of DT, or lambda of L1 reg. Remember, there are several other hyper parameters.

Report the performance of the test-data. Create a similar visualization with 9 methods now, with your additional 4 methods. The plot shows the mean and variance, FYI. Use a suitable visualization method to get them. You may wonder; to compute variance, you need more than 2 samples. Right. Repeat this task 5 times to get the mean and variance.

Importing Libraries

```
[1]: import pandas as pd # type: ignore
import numpy as np # type: ignore
import seaborn as sns # type: ignore
import matplotlib.pyplot as plt # type: ignore

from sklearn.model_selection import train_test_split, GridSearchCV # type: ignore
from sklearn import tree # type: ignore
from sklearn.tree import DecisionTreeClassifier # type: ignore
from sklearn.metrics import accuracy_score, precision_score # type: ignore
```

```
import warnings
warnings.filterwarnings("ignore")
```

Data Processing

```
[2]: file_path = '/Users/nandhakishorecs/Documents/IITM/Jul_Nov_2024/
      ↪DA5401_Data_Analytics_Lab/Assignment5/nursery/nursery.data'
df = pd.read_csv(file_path, names = ['parents', 'has_nurs', 'form', 'children', 'housing', 'finance', 'social', 'final_evaluation'])
```

```
[3]: df.describe()
```

```
[3]:
```

	parents	has_nurs	form	children	housing	finance	social \
count	12960	12960	12960	12960	12960	12960	12960
unique	3	5	4	4	3	2	3
top	usual	proper	complete	1	convenient	convenient	nonprob
freq	4320	2592	3240	3240	4320	6480	4320

	health	final_evaluation
count	12960	12960
unique	3	5
top	recommended	not_recom
freq	4320	4320

```
[4]: df['final_evaluation'].unique()
```

```
[4]: array(['recommend', 'priority', 'not_recom', 'very_recom', 'spec_prior'],
      dtype=object)
```

Collapsing 5 label dataset to a 3 label dataset

```
[5]: # OLD: 0 - recommend, 1 - priority, 2 - not_recom, 3 - very_recom, 4 - spec_prior
      ↪spec_prior
# NEW: 0 - [spec_prior, recommend, very_recom], 1 - not_recom, 2 - priority

new_classes = { 'recommend':'recommended', 'priority':'priority' , 'not_recom':
      ↪'not_recom', 'very_recom':'recommended', 'spec_prior':'recommended' }
df['final_evaluation'] = df['final_evaluation'].map(new_classes)
df
```

```
[5]:
```

	parents	has_nurs	form	children	housing	finance \
0	usual	proper	complete	1	convenient	convenient
1	usual	proper	complete	1	convenient	convenient
2	usual	proper	complete	1	convenient	convenient
3	usual	proper	complete	1	convenient	convenient
4	usual	proper	complete	1	convenient	convenient
...

12955	great_pret	very_crit	foster	more	critical	inconv
12956	great_pret	very_crit	foster	more	critical	inconv
12957	great_pret	very_crit	foster	more	critical	inconv
12958	great_pret	very_crit	foster	more	critical	inconv
12959	great_pret	very_crit	foster	more	critical	inconv

	social	health	final	evaluation
0	nonprob	recommended		recommended
1	nonprob	priority		priority
2	nonprob	not_recom		not_recom
3	slightly_prob	recommended		recommended
4	slightly_prob	priority		priority
...
12955	slightly_prob	priority		recommended
12956	slightly_prob	not_recom		not_recom
12957	problematic	recommended		recommended
12958	problematic	priority		recommended
12959	problematic	not_recom		not_recom

[12960 rows x 9 columns]

Checking for missing values and Nan

```
[6]: print(df['final evaluation'].unique())
df.isnull().sum()
```

```
['recommended' 'priority' 'not_recom']
```

```
[6]: parents          0
has_nurs             0
form                 0
children             0
housing              0
finance              0
social               0
health               0
final evaluation     0
dtype: int64
```

Performing Hyper-parameter tuning for the models:

1. Decision Tree Classifier
2. Decision Tree Classifier - Using OneHot Encoding
3. Logistic Regression with L1 Regularisation
4. K Nearest Neighbour Classifier

```
[7]: from sklearn.preprocessing import LabelEncoder # type: ignore
```

```
df_label_encoded = df.apply(LabelEncoder().fit_transform)
df_label_encoded.columns
```

```
[7]: Index(['parents', 'has_nurs', 'form', 'children', 'housing', 'finance',
          'social', 'health', 'final evaluation'],
          dtype='object')
```

```
[8]: from sklearn.preprocessing import OneHotEncoder # type: ignore
      encoder = OneHotEncoder(sparse_output=False)

      encoded_table = encoder.fit_transform(df.drop('final evaluation', axis = 1))

      df_onehot_encoded = pd.DataFrame(encoded_table, columns=encoder.
      ↪get_feature_names_out())
      df_onehot_encoded['final evaluation'] = df_label_encoded['final evaluation']
```

```
[ ]: # performing the experiment for 5 times to get the mean and variance of ↪
      ↪accuracy values

      from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression # type: ignore
      from sklearn.neighbors import KNeighborsClassifier # type: ignore

      # creating separate dictionaries

      accuracy_results = {
          'Decision_Tree_LE': [],
          'Decision_Tree_OHE': [],
          'Logistic_Regression_L1_OHE': [],
          'KNN_OHE': []
      }
      precision_results = {
          'Decision_Tree_LE': [],
          'Decision_Tree_OHE': [],
          'Logistic_Regression_L1_OHE': [],
          'KNN_OHE': []
      }

      for i in range(5):

          #----- Decision Tree - Label Encoding ↪
          ↪-----

          label_train_x, label_test_x, label_train_y, label_test_y = train_test_split(
              df_label_encoded.drop('final evaluation', axis = 1),
              df_label_encoded['final evaluation'],
              test_size = 0.2
          )
```

```

param_grid = {
    'max_depth' : [1, 2, 3, 4, None],
    'min_samples_split' : [2, 5, 10],
    'min_samples_leaf' : [1, 2, 4, 8, 16],
}
tree_clf = DecisionTreeClassifier()
grid_search = GridSearchCV(
    estimator = tree_clf,
    param_grid = param_grid,
    cv = 5,
    n_jobs = -1,
    verbose = 2,
    scoring = 'neg_mean_squared_error'
)
grid_search.fit(label_train_x, label_train_y)

best_tree_clf = grid_search.best_estimator_
pred_y = best_tree_clf.predict(label_test_x)
dt_test_acc = accuracy_score(label_test_y, pred_y)
dt_test_precision = precision_score(label_test_y, pred_y, average = 'micro')
accuracy_results['Decision_Tree_LE'].append(dt_test_acc)
precision_results['Decision_Tree_LE'].append(dt_test_precision)

#----- Decision Tree - One Hot Encoding
↳ -----

one_hot_train_x, one_hot_test_x, one_hot_train_y, one_hot_test_y =
↳ train_test_split(
    df_onehot_encoded.drop('final evaluation', axis = 1),
    df_label_encoded['final evaluation'],
    test_size = 0.2
)

param_grid = {
    'max_depth' : [1, 2, 3, 4, None],
    'min_samples_split' : [2, 5, 10],
    'min_samples_leaf' : [1, 2, 4, 8, 16]
}
tree_clf = DecisionTreeClassifier()
grid_search = GridSearchCV(
    estimator = tree_clf,
    param_grid = param_grid,
    cv = 5,
    n_jobs = -1,
    verbose = 2,
    scoring = 'neg_mean_squared_error'
)

```

```

)
grid_search.fit(one_hot_train_x, one_hot_train_y)

best_tree_clf = grid_search.best_estimator_
pred_y = best_tree_clf.predict(one_hot_test_x)
dt_ohe_test_acc = accuracy_score(one_hot_test_y, pred_y)
df_ohe_test_precision = precision_score(one_hot_test_y, pred_y, average =
↳ 'micro')

accuracy_results['Decision_Tree_OHE'].append(dt_ohe_test_acc)
precision_results['Decision_Tree_OHE'].append(df_ohe_test_precision)

#----- Lasso Logistic Regression - One Hot Encoding
↳ -----

parameter_grid = [{
    'C': [10000, 1000, 100, 10, 1, 0.1, 0.01, 0.001, 0.0001],
    # 'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    # 'solver': ['lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'],
    'max_iter': [1, 2, 4, 8, 16, 32]
}]

clf = LogisticRegression(penalty='l1', solver='liblinear')
grid_search = GridSearchCV(
    estimator = clf,
    param_grid = parameter_grid,
    cv = 5,
    n_jobs = -1,
    verbose = True,
    scoring = 'neg_mean_squared_error'
)
grid_search.fit(one_hot_train_x, one_hot_train_y)

clf = grid_search.best_estimator_
pred_y = clf.predict(one_hot_test_x)
lr_test_acc = accuracy_score(one_hot_test_y, pred_y)
lr_test_precision = precision_score(one_hot_test_y, pred_y, average =
↳ 'micro')

accuracy_results['Logistic_Regression_L1_OHE'].append(lr_test_acc)
precision_results['Logistic_Regression_L1_OHE'].append(lr_test_precision)

#----- K Nearest Neighbor Classifier - One Hot
↳ Encoding -----

parameter_grid = {

```

```

        'n_neighbors': [ 3, 5, 7, 9],
        'metric': ['euclidean', 'manhattan', 'minkowski'],
        'weights': ['uniform', 'distance'],
        'algorithm' : ['auto', 'kd_tree', 'brute', 'ball_tree']
    }

    knn = KNeighborsClassifier()
    grid_search = GridSearchCV(
        estimator = knn,
        param_grid = parameter_grid,
        cv = 5,
        n_jobs = -1,
        verbose = True,
    )
    grid_search.fit(one_hot_train_x, one_hot_train_y)

    clf = grid_search.best_estimator_
    pred_y = clf.predict(one_hot_test_x)
    knn_test_acc = accuracy_score(one_hot_test_y, pred_y)
    knn_test_precision = precision_score(one_hot_test_y, pred_y, average =
    ↪ 'micro')

    accuracy_results['KNN_OHE'].append(lr_test_precision)
    precision_results['KNN_OHE'].append(lr_test_precision)

```

Helper function to get mean and variance to get confidence interval to plot the error plots

```

[10]: import scipy # type: ignore
from scipy.stats import t, sem # type: ignore
\
# Referece: https://sphweb.bumc.bu.edu/otlt/MPH-Modules/PH717-QuantCore/
    ↪ PH717-Module6-RandomError/PH717-Module6-RandomError11.html

def mean_confidence_interval(data, confidence_interval = 0.95 ):
    mean = np.mean(data)
    sem = scipy.stats.sem(data)
    interval = scipy.stats.t.interval(confidence_interval, len(data)-1, loc =
    ↪ mean, scale =sem)
    return mean , interval

```

```

[11]: models = list(accuracy_results.keys())
type(models)
print(models)

```

```

['Decision_Tree_LE', 'Decision_Tree_OHE', 'Logistic_Regression_L1_OHE',
'KNN_OHE']

```

Plotting the Accuracy values for our models

```

[12]: acc_mean_values = []
      acc_upper_bounds = []
      acc_lower_bounds = []

      for model in models:
          mean, ci = mean_confidence_interval(accuracy_results[model])
          mean = np.array(mean) * 100
          ci = np.array(ci) * 100

          lower_bound = mean - ci[0]
          upper_bound = - mean + ci[1]

          acc_mean_values.append(mean)
          acc_upper_bounds.append(upper_bound)
          acc_lower_bounds.append(lower_bound)

      ## Getting the accuracy values of other models from the webpage for plotting
      ↳ error plot

      # Getting Model Names
      models.append('Xgboost Classification')
      models.append('SVM')
      models.append('Random Forest Classification')
      models.append('Neural Network Classification')
      models.append('Logistic Regression')

      print(models)

      # Adding mean values from the webpage
      acc_mean_values.append(99.969)
      acc_mean_values.append(99.198)
      acc_mean_values.append(98.025)
      acc_mean_values.append(100.000)
      acc_mean_values.append(92.253)

      # Adding boundary values from the webpage

      acc_lower_bounds.append(acc_mean_values[4] - 99.907)
      acc_lower_bounds.append(acc_mean_values[5] - 98.889)
      acc_lower_bounds.append(acc_mean_values[6] - 97.531)
      acc_lower_bounds.append(acc_mean_values[7] - 100.000)
      acc_lower_bounds.append(acc_mean_values[8] - 91.327)

      acc_upper_bounds.append(100.000 - acc_mean_values[4])
      acc_upper_bounds.append(99.475 - acc_mean_values[5])

```



```

acc_upper_bounds.append(98.488 - acc_mean_values[6])
acc_upper_bounds.append(100.000 - acc_mean_values[7])
acc_upper_bounds.append(93.117 - acc_mean_values[8])

print(acc_lower_bounds)

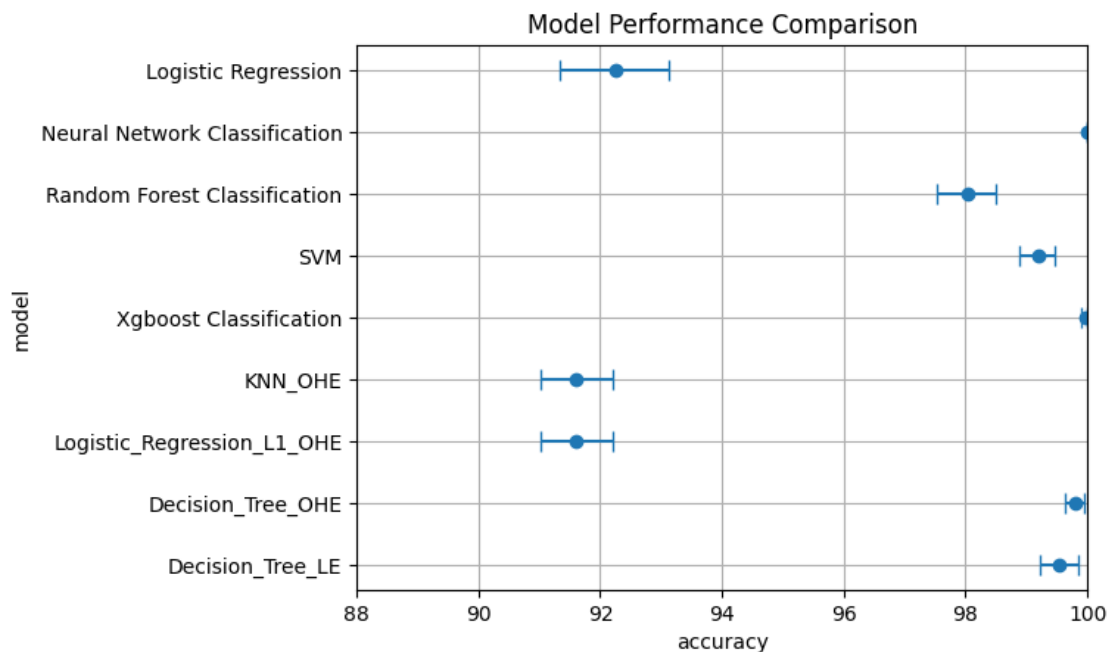
plt.errorbar(acc_mean_values, models, xerr=[acc_lower_bounds,
↪acc_upper_bounds], fmt='o', capsize=5)
plt.xlim(88, 100)
plt.xlabel('accuracy')
plt.ylabel('model')
plt.title('Model Performance Comparison')
plt.grid()

```

```

['Decision_Tree_LE', 'Decision_Tree_OHE', 'Logistic_Regression_L1_OHE',
'KNN_OHE', 'Xgboost Classification', 'SVM', 'Random Forest Classification',
'Neural Network Classification', 'Logistic Regression']
[0.32489872546237564, 0.16455462587948944, 0.5911790611771295,
0.5911790611771295, 0.06199999999999761, 0.3089999999999975, 0.4939999999999998,
0.0, 0.92600000000000019]

```



```

[13]: # Tried Precision value plotting in the hope of replicating the results from
↪webpage
'''

```

```

pre_mean_values = []
pre_upper_bounds = []
pre_lower_bounds = []

for model in model_names:
    mean, ci = mean_confidence_interval(precision_results[model])
    mean = np.array(mean) * 100
    ci = np.array(ci) * 100

    lower_bound = mean - ci[0]
    upper_bound = - mean + ci[1]

    pre_mean_values.append(mean)
    pre_upper_bounds.append(upper_bound)
    pre_lower_bounds.append(lower_bound)

plt.errorbar(pre_mean_values, model_names, xerr=[acc_lower_bounds,
↪acc_upper_bounds], fmt='o', capsize=5)
plt.xlabel('precision')
plt.ylabel('model')
plt.title('Model Performance Comparison')

plt.grid()
'''

```

```

[13]: "\n    pre_mean_values = []\n    pre_upper_bounds = []\n    pre_lower_bounds =
[]\n\n    for model in model_names: \n        mean, ci =
mean_confidence_interval(precision_results[model])\n        mean =
np.array(mean) * 100 \n        ci = np.array(ci) * 100 \n\n        lower_bound =
mean - ci[0]\n        upper_bound = - mean + ci[1]\n\n
pre_mean_values.append(mean)\n        pre_upper_bounds.append(upper_bound)\n
pre_lower_bounds.append(lower_bound)\n\n    plt.errorbar(pre_mean_values,
model_names, xerr=[acc_lower_bounds, acc_upper_bounds], fmt='o', capsize=5)\n
plt.xlabel('precision')\n    plt.ylabel('model')\n    plt.title('Model
Performance Comparison')\n\n    plt.grid() \n"

```

Observations

1. From the data from dataset's webpage, clearly Neural Networks and Xgboost are overfitting.
2. For Logistic Regression with L1 Regularisation, even for larget regularisation constant, the model performs worse than the baseline logistic regression model
3. Decision Trees - both with label encoding and one hot encoding, the model is performing well with an accuracy of 99% and above
4. Kor KNN, the performance is on par with the Logistic regression with L1 regularisation

Task 2 [10 Points] (From Question)

You may notice that the shape of logistic regression decision boundary and a sigmoid are a look-

alike. We know that range of sigmoid is 0 to 1, which means, we can use sigmoid only when outputs are unipolar. Here are some simple extensions, we may try

1. Construct a `bipolar_sigmoid(x)` using unipolar sigmoid.
2. A popular bipolar normalizer is $\tanh(x)$. Compare the response of $\tanh(x)$ vs your `bipolar_sigmoid(x)`.
3. Parameterize it as `bipolar_sigmoid(ax)`, $\tanh(ax)$; You may plot the shapes of the response at different values of 'a' in [-5, -1, -0.1, -0.01, .001, .01, .1, 1, 5].
4. Now comes the interesting part. Can you evaluate the linear range of 'x' for each value of 'a' in `bipolar_sigmoid(ax)`? Usually, when 'a' is small, the linearity range is high.

Creating a class 'Sigmoid' with a family of sigmoid functions including unipolar sigmoid, bipolar sigmoid and tanh

```
[14]: from math import exp

class sigmoids:
    def unipolar_sigmoid(self, x:float, a:float = 1) -> float:
        return 1 / (1 + exp(-1 * a * x))


    def bipolar_sigmoid(self, x:float, a:float = 1) -> float:
        # return (1 - exp(-1 * a * x))/(1 + exp(-1 * a * x))
        return (2 * self.unipolar_sigmoid(x, a)) - 1

    def tanh(self, x:float, a:float = 1) -> float:
        return (exp(a*x) - exp(-1*a*x))/(exp(a*x) + exp(-1*a*x))
```

Plotting sigmoid function(s) with a = 1

```
[15]: sigmoids = sigmoids()
x = np.linspace(-10, 10, 101)
y_unipolar_sigmoid = []
y_bipolar_sigmoid = []
y_tanh = []
for i in range(0, len(x)):
    y_unipolar_sigmoid.append(sigmoids.unipolar_sigmoid(x[i]))
    y_bipolar_sigmoid.append(sigmoids.bipolar_sigmoid(x[i]))
    y_tanh.append(sigmoids.tanh(x[i]))

plt.plot(x, y_unipolar_sigmoid)
plt.plot(x, y_bipolar_sigmoid)
plt.plot(x, y_tanh)
plt.title('Unipolar Sigmoid , Bipolar Sigmoid, Hyperbolic Tangent')
plt.legend(['Unipolar', 'Bipolar', 'Hyperbolic Tangent'], loc = 'lower right')
plt.grid()
plt.show()
```



Plotting sigmoid function(s) with different values of a

```
[16]: x = np.linspace(-10, 10, 11)

y_unipolar_sigmoid = {}
y_bipolar_sigmoid = {}
y_tanh = {}

a = [-5, -1, -.1, -.01, .001, .01, .1, 1, 5]

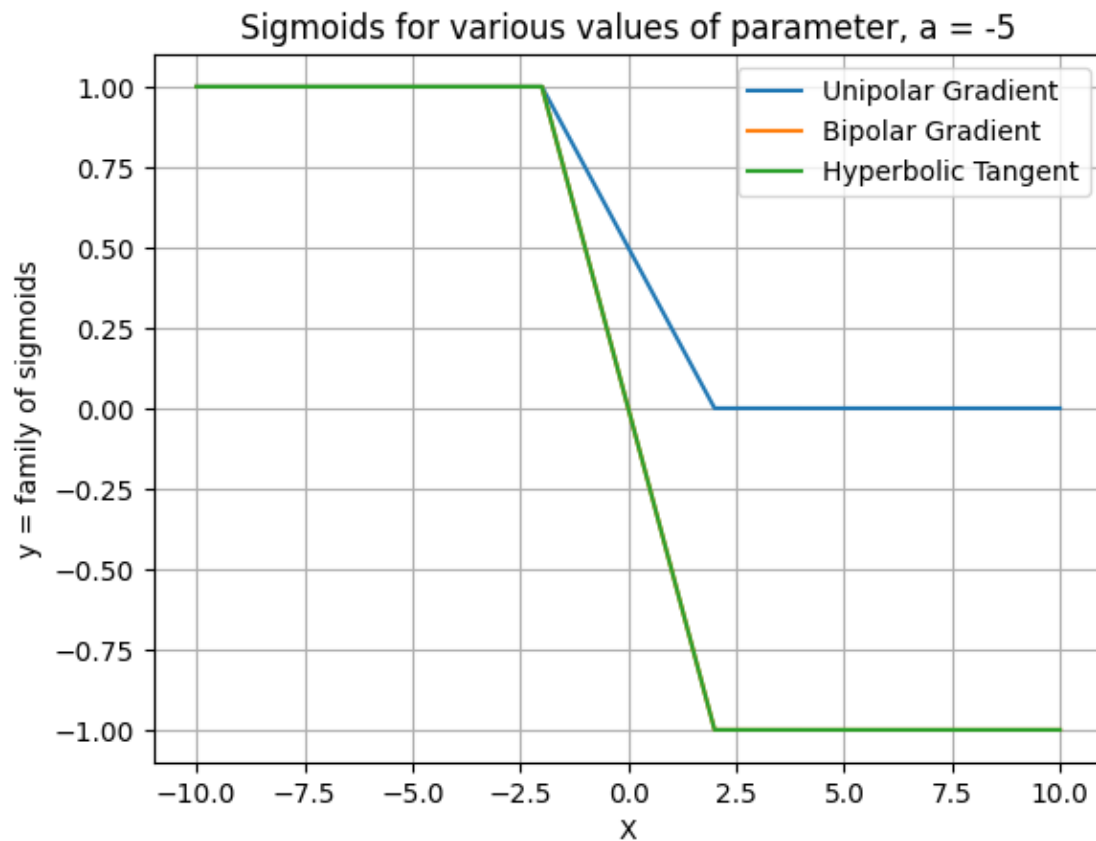
for i in a:
    y_unipolar_sigmoid[i] = []
    y_bipolar_sigmoid[i] = []
    y_tanh[i] = []


for j in a:
    for i in x:
        y_unipolar_sigmoid[j].append(sigmoids.unipolar_sigmoid(i, j))
```

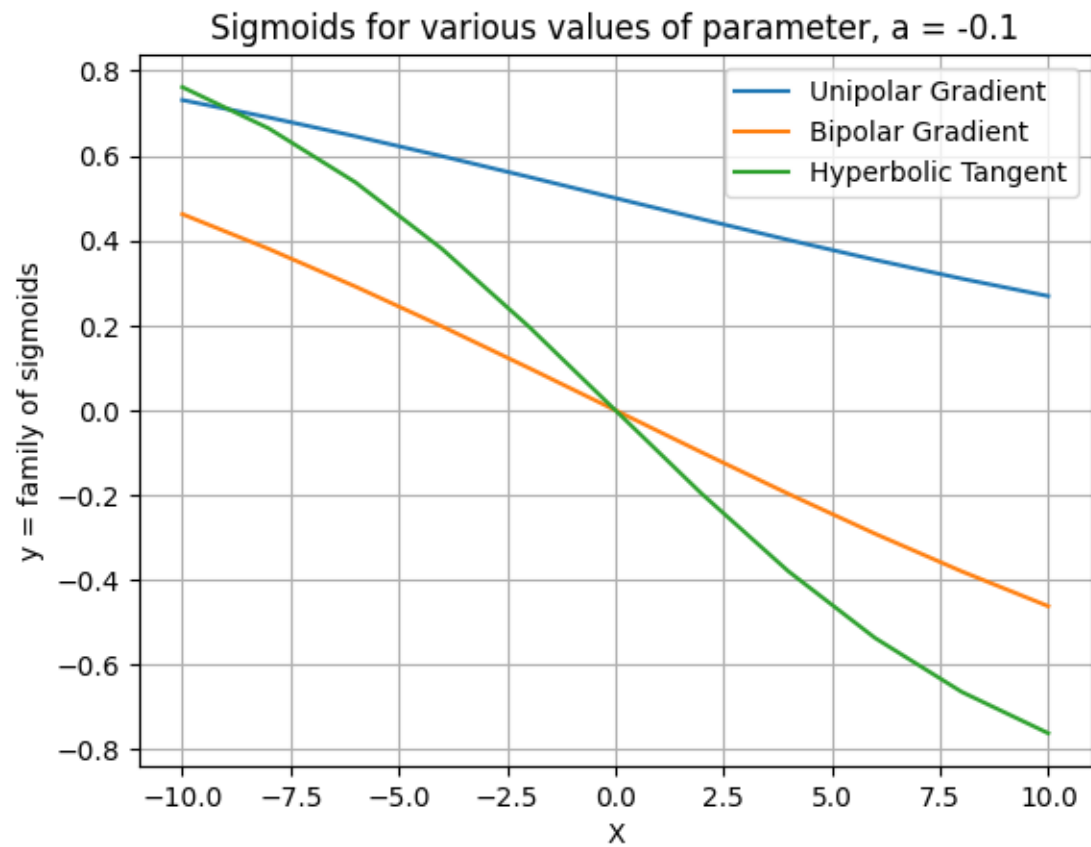
```
y_bipolar_sigmoid[j].append(sigmoids.bipolar_sigmoid(i, j))
```

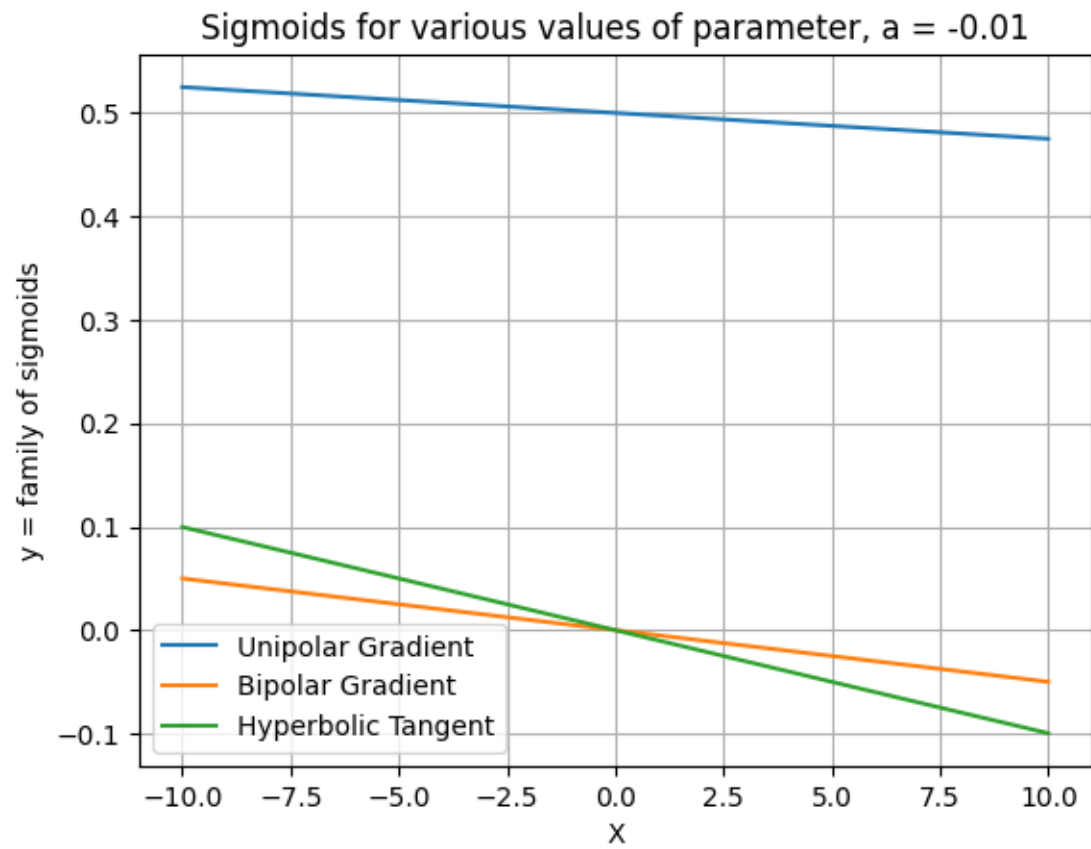
```
y_tanh[j].append(sigmoids.tanh(i, j))
```

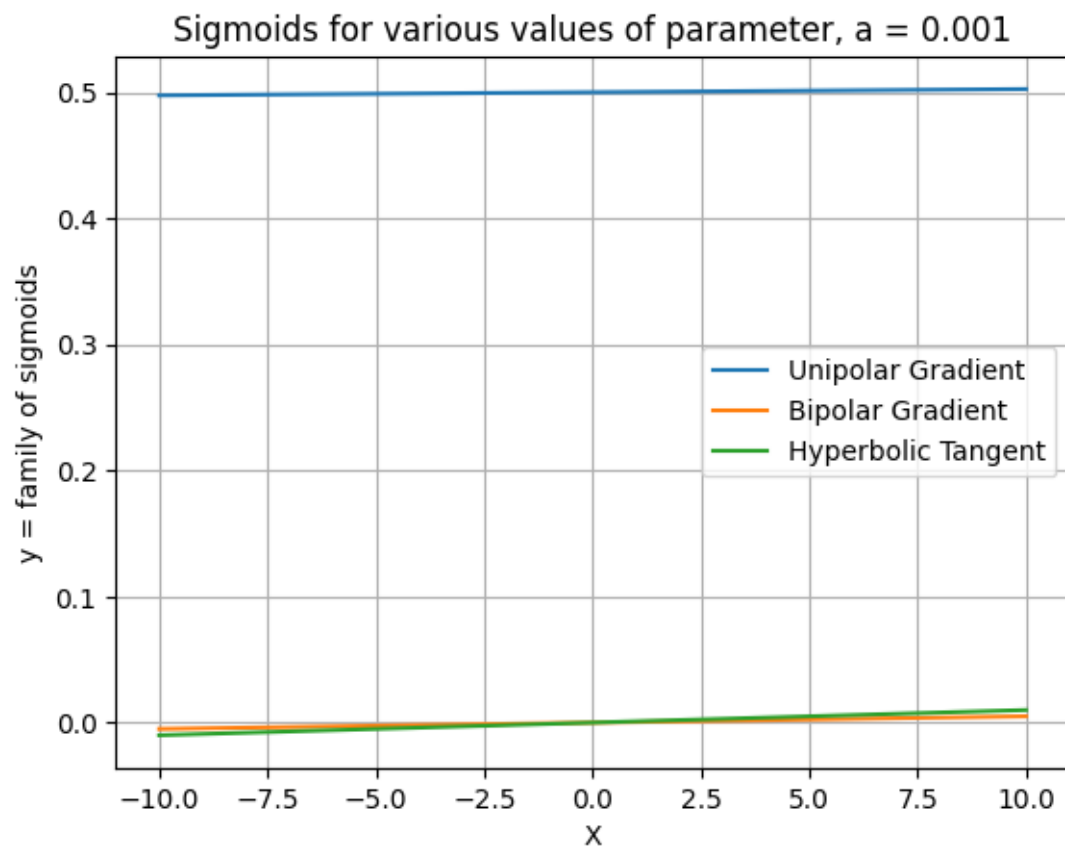
```
[17]: for parameter in a:
    plt.plot(x, y_unipolar_sigmoid[parameter])
    plt.plot(x, y_bipolar_sigmoid[parameter])
    plt.plot(x, y_tanh[parameter])
    plt.title(f'Sigmoids for various values of parameter, a = {parameter}')
    plt.xlabel('X')
    plt.ylabel('y = family of sigmoids')
    plt.legend(['Unipolar Gradient', 'Bipolar Gradient', 'Hyperbolic Tangent'], loc = 'best')
    plt.grid()
    plt.show()
```

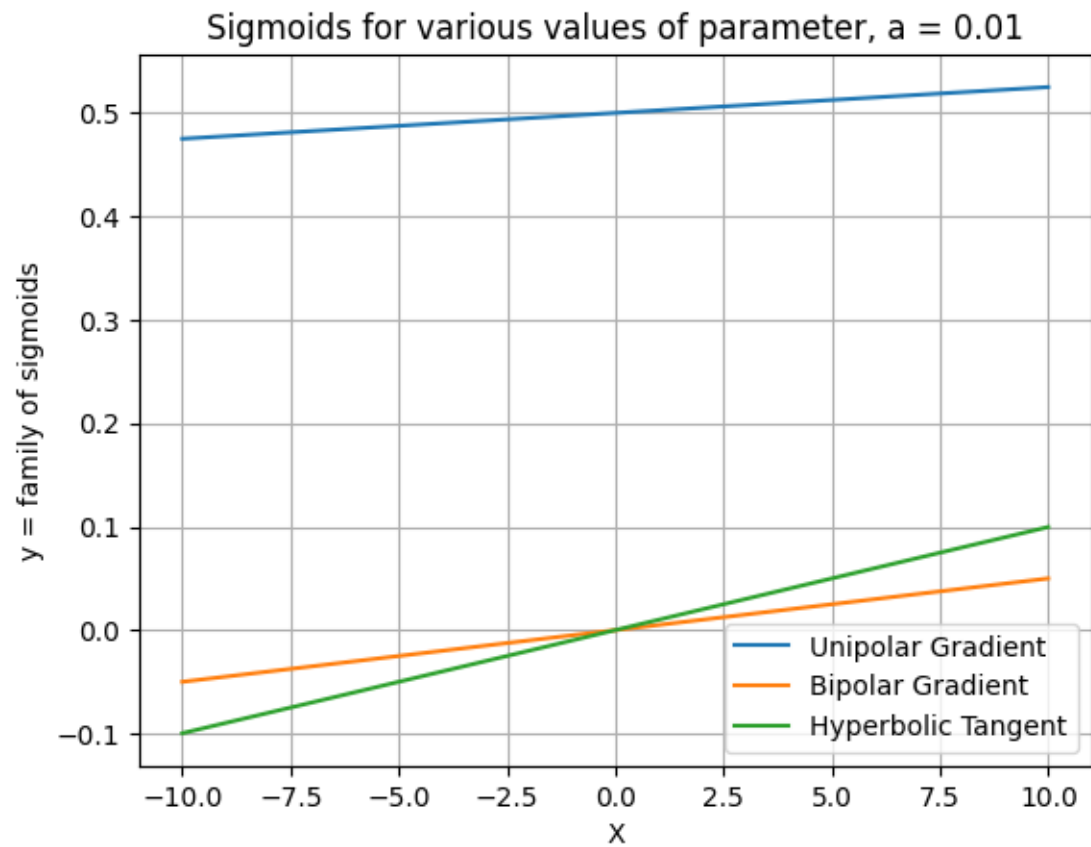





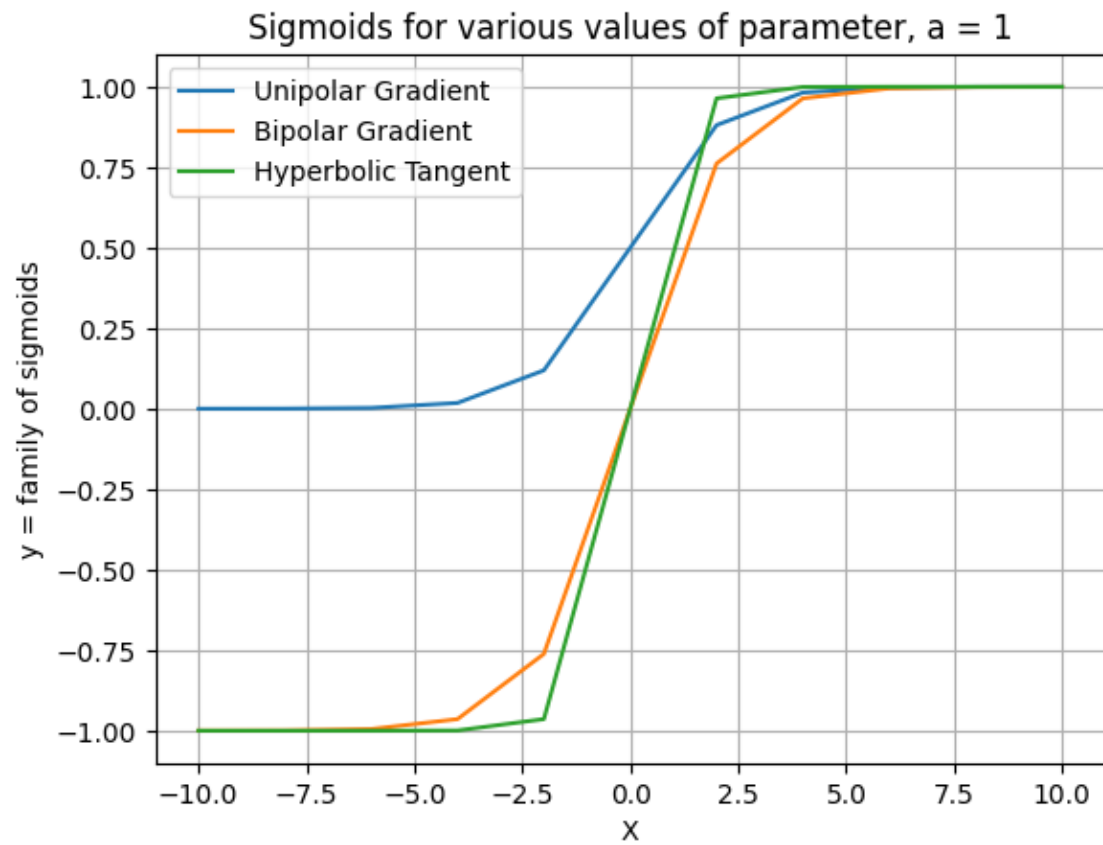


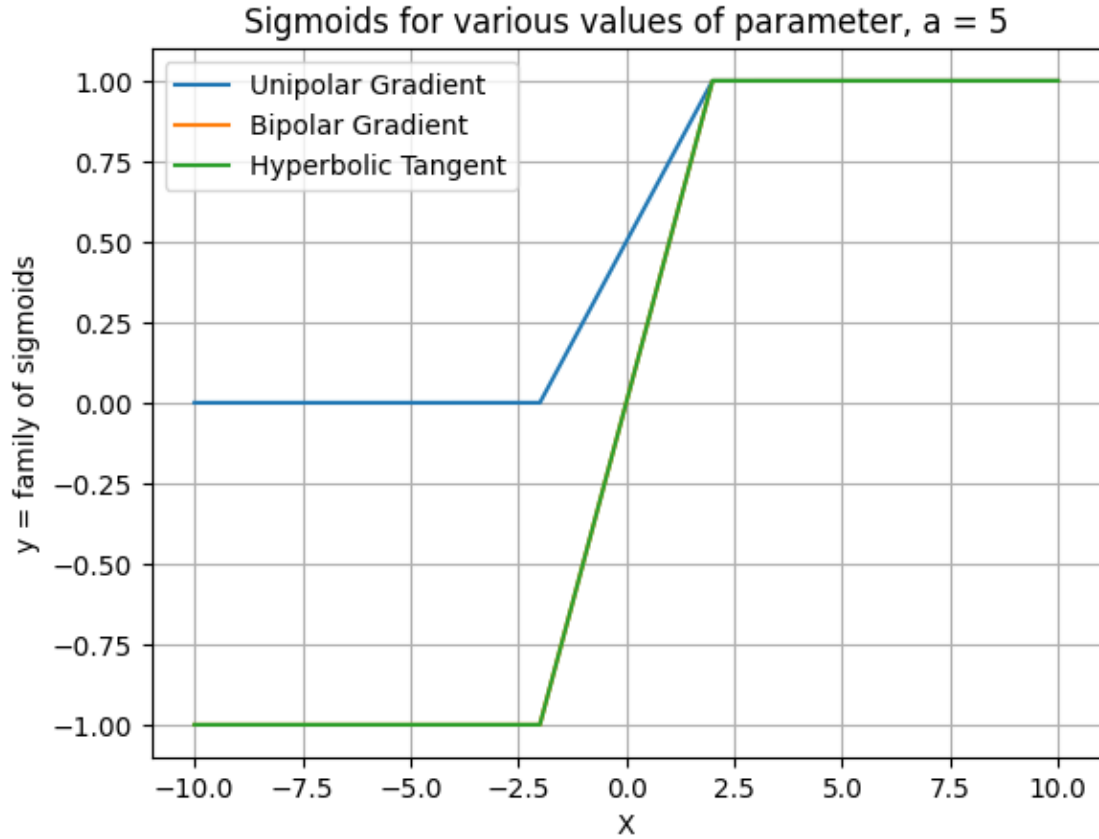












Observations

1. For values of 'a' which are far away from zero, the sigmoids (unipolar, bipolar and hyperbolic tangent) tend to be a curve and for 'a' = 0, the sigmoids are straight lines.
2. In scenarios where the sigmoid is a curve (i.e.) the values of parameter 'a' are far from zero, the larger the value of 'a', the steeper the curve climbs from negative (y) to positive (y).
3. For values of 'a' very close to zero or zero, the sigmoid functions act as straight lines. This can be understood by taking the gradient of the sigmoid function for the given particular value of 'a'. Note that, for sigmoid, the derivative is $\sigma(\sigma - 1)$, thus when the value of 'a' is very close to zero, the gradient of the curve increases.